Do bridge-monitoring error bars survive an unseen season? Physics-guided, conformally calibrated temperature compensation on the KW51 railway bridge

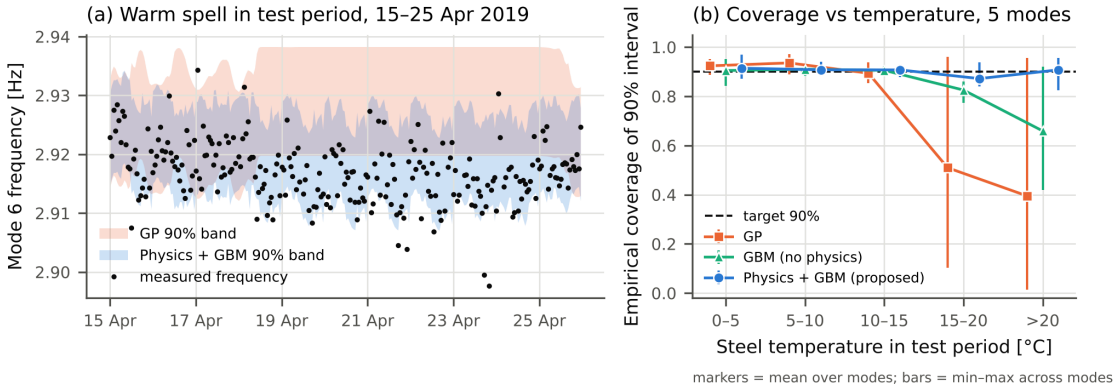
Ethan Li · High school student, Austin, TX · poggies55555@gmail.com · updated 10 November 2026
Code and results: github.com/apersonli55555/fall-detector (folder research/bridge-uncertainty)

PROBLEM

Bridges are monitored by tracking natural frequencies, which fall when damage lowers stiffness. Temperature shifts them by a similar amount, so alarms depend on a temperature-compensated uncertainty band. Conformal prediction guarantees coverage only if past and future data are exchangeable; seasonal data are not. **How much does coverage degrade when a model must predict a season it has not seen, and does a physics prior help?**

DATA AND METHOD

I used the public KW51 dataset (Maes & Lombaert, Zenodo): hourly natural frequencies of a steel railway bridge in Leuven, Belgium, Oct 2018–Jan 2020, retrofitted in May–Sep 2019. No new experiments were run. I trained on Nov–Jan, calibrated on Feb, and tested on Mar–mid-May, when temperatures reached 26 °C against a training maximum of 13.6 °C. The proposed model is a 4-parameter thermo-elastic curve, $f = f_0\sqrt{(1+a\Delta T)}$ plus a cold-stiffening term, with a small gradient-boosted model fitted to its residuals. It was compared with linear, cubic, Gaussian-process and physics-free boosted-tree baselines, all wrapped in the same split-conformal procedure. Five vibration modes were analysed, with 9 models including ablations, 5 designed folds and 22 rolling windows.



Key figure. (a) Mode 6 during a warm spell in the test period: the Gaussian-process 90 % band drifts away from the measurements, while the physics-guided band follows them. (b) Empirical coverage of nominal 90 % intervals by test temperature, averaged over 5 modes (bars: min–max across modes). Training never exceeded 13.6 °C.

RESULTS SO FAR

- With a random (exchangeable) split, all 9 models were calibrated (0.90 ± 0.01), as theory predicts.
- On test hours warmer than anything in training (22 rolling windows × 5 modes), the physics-guided model held a median of **0.91** coverage. In paired tests it beat all 8 single-regime and non-parametric models: +0.08 [0.05, 0.12] over boosted trees, +0.36 over a GP (adjusted $p \le 0.002$).
- **Why, and a correction (10 Nov):** what matters is separating the freezing regime from the warm-weather trend. Without the cold term the fitted slope is up to 2× too steep and coverage drops from 0.91 to 0.83. But a *physics-free* straight line with the same kink matches my model (0.896 vs 0.895; difference −0.002 [−0.006, 0.002]). The gain comes from the two-regime shape, not from the physics formula.
- **Negative result:** tested on colder hours than training, every model, mine included, failed (median coverage 0.50–0.62), and after the retrofit mine was significantly worse than simple baselines. Physics cannot extrapolate into a mechanism the data never showed.
- Frozen bands flagged the 2019 retrofit on every post-retrofit day for 3 modes (0–1.3 % false-alarm days).

LIMITATIONS AND NEXT STEPS

This is one bridge and one season, so the rolling windows overlap. The retrofit was a large stiffening, not damage, so a minimum-detectable-shift study is needed. Next: a cross-season prior for the cold term, a second structure (LUMO), and a physics-informed neural-network residual. **I would value feedback on the physical cause of the cold stiffening and on evaluation design.**

All numbers are reproducible from the project repository. The code was developed with the help of an AI coding assistant. Prior work on KW51 (Maes et al. 2022; Wittenberg et al. 2025; Neumann et al. 2025–26) covers environmental compensation and monitoring. In my search I did not find conformal coverage evaluated under seasonal shift for this dataset.