How Much Is a Real Fall Worth? Real, Synthetic and Staged Training Data for Video Fall Detection, and Adapting with a Handful of Real Clips

Anonymous submission

Abstract

Video fall detectors are commonly trained and evaluated on falls staged by actors in a laboratory, where they score near perfectly. We ask which training data actually prepares a detector for real falls. We build **FallOrNot**, a real-world test set from the OOPS! dataset of in-the-wild accident videos, sampled independently of any model and labelled for whether a person falls, so that *other accidents* act as false-alarm negatives. Using frozen CLIP and VideoMAE features with linear probes, and end-to-end fine-tuned VideoMAE, we compare staged (UR Fall), synthetic (OmniFall OF-Syn) and real (OOPS!) training data, and test on FallOrNot, on OmniFall's larger in-the-wild test split (OF-ItW, labelled by an independent team) and on an independent real-world source, HMDB51. Staged-trained probes that score 1.00 in-domain reach 0.76 AUROC on FallOrNot. On FallOrNot, 131 real training falls lift the best probe to 0.90, and about 30 real falls outperform 2,136 synthetic ones, whose benefit plateaus near 0.82; OF-ItW confirms the ordering (staged 0.82, synthetic 0.87, real 0.93–0.95). On HMDB51, however, the ranking flips: synthetic training is best and real training is no better than staged. The flip holds for fine-tuned models, averaged over three seeds. Training on synthetic and real data together never falls more than 0.03 AUROC below the best source on any test, whereas every single source falls at least 0.09 below it on some test. For the common case where only a handful of real falls is available, we test an *anchored probe*: a probe trained on the real clips whose weights are penalised for leaving the synthetic-trained probe's weights. Under a success rule fixed before any result, it stays within 0.005 AUROC of the better of the real-only and synthetic-only probes at every budget from 4 to 125 real falls, and is higher than both on average, on FallOrNot and OF-ItW (with 4 real falls: 0.81 vs. 0.69 real-only on FallOrNot); with the full real pool it gives no gain over real-only training. Two other tests specified in advance (fine-tuning beating frozen probes, and source-balanced mixing beating plain mixing) were not met. Real training clips from compilations absent from the test set give the same result as random subsets of equal size, so shared compilations do not explain the real-data advantage. A second annotator agreed moderately with the first ($\kappa = 0.50$); re-grading with their labels preserves the main findings but weakens the in-domain real-over-synthetic advantage. Conclusions about training data drawn from a single real-world benchmark do not generalise; fall detectors should be evaluated on several.

1 Introduction

Falls are the second-leading cause of unintentional injury death worldwide, and about 71% of fall deaths occur in people aged 60 or older [18]; long periods on the floor after a fall are associated with worse outcomes [5]. Camera-based fall detection promises automatic alerts without wearables, and published detectors often report near-perfect accuracy. Most of that evidence comes from staged datasets in which young actors fall onto mattresses in front of fixed cameras [11, 3].

Real falls are different: they are sudden, partly occluded, filmed from arbitrary viewpoints, and embedded among other abrupt events such as collisions, dropped objects and people catching themselves. OmniFall [13] recently showed a large drop from staged to in-the-wild falls on fall-containing OOPS! [4] videos, and argued that synthetic training videos generalise better than staged ones. This leaves practical questions open for anyone building a detector: *how much real data is needed, is real data worth more than synthetic data, does the answer change with the test set, and how often do detectors raise alarms for accidents that are not falls?*

Contributions.

1. **FallOrNot**, a real-world test set sampled from *all* annotated OOPS! accidents independently of any model, so that non-fall accidents are hard false-alarm negatives, plus pre-accident clips of ordinary activity.
2. A controlled comparison of **staged, synthetic and real training data** with identical encoders, probes and clips, including an equal-budget curve and a learning curve over the number of real falls.
3. Evidence that **rankings flip between real-world test sets**: real data wins on FallOrNot, synthetic data wins on an independent source (HMDB51), for both frozen probes and fine-tuned models; synthetic + real has the smallest worst-case shortfall.
4. An **anchored probe** for adapting a synthetic-trained detector with a handful of real clips, tested under a success rule **specified in advance** and met on two test sets (Section 5.6).

5. Two further tests specified in advance (fine-tuning and source-balanced mixing; neither met), confirmation of the in-domain ordering on OmniFall's **OF-ItW** test split, a compilation leakage check, a second-annotator agreement study, and a re-analysis with the second annotator's labels.

2 Related work

Staged fall datasets. UR Fall [11] records 30 falls from a side and a ceiling camera and 40 activities of daily living; Le2i [3] and the Multiple Cameras Fall Dataset [1] are other widely used staged sets. Detectors evaluated on such data routinely approach ceiling performance.

Simulated vs. real falls. For wearable accelerometers, algorithms tuned on simulated falls lose accuracy on real-world falls [2], and the FARSEEING repository collects body-worn sensor signals from real falls [8]. We study the analogous question for video.

In-the-wild video. OOPS! [4] contains about 20k web videos of unintentional actions with annotated transition times. **OmniFall** [13], the closest work, unifies eight staged datasets, adds 12k synthetic videos (OF-Syn) and an in-the-wild test set of 818 fall-containing OOPS! videos (OF-ItW), shows a large staged-to-wild drop, finds synthetic training better than staged, and reports an upper-bound model trained with in-the-wild data. We do not claim to be first to show the staged-to-wild gap. Our work differs in (i) negatives drawn from all accidents by a model-independent sample, (ii) the question of how much real data is needed and whether it beats synthetic data at equal budget, and (iii) testing every conclusion on a second, independent real-world source.

Encoders. We use CLIP [15] through OpenCLIP [14] with DataComp weights [6], and VideoMAE [16] fine-tuned on Kinetics-400 [7].

3 The FallOrNot test set

Source and clips. We use the 10,961 OOPS! videos whose transition time t was found by at least two annotators (median used). Each video yields an *event* clip $[t-1, t+1]$ s and, when $t \geq 3.5$ s, a *pre-event* clip $[t-3.5, t-1.5]$ s of ordinary activity. Every clip, in every dataset in this paper, is 2 s long and sampled at 16 frames.

Model-independent test sample. A video enters the test sample if the MD5 hash of its name falls below a fixed threshold (about 600 videos; 422 with usable clips). Selection depends on no model, so the benchmark cannot favour any detector. Clips were labelled in a fixed random order, so any labelled prefix is itself a random sample.

Labels. Following the WHO definition, a fall is an event in which a person inadvertently comes to rest on the ground, floor or another lower level [18]. The annotator answered *fall*, *not a fall* (including other accidents, near-falls and animals) or *unsure* (excluded). Positives are event clips labelled *fall*; negatives are event clips labelled *not a fall* (hard negatives) and the pre-event clips of labelled test videos (easy negatives).

Size. 165 test videos are labelled (255 clips, 102 falls). A separate training pool, the 3,000 event clips with the highest zero-shot CLIP fall score (never overlapping the test sample), provides 268 labelled real training clips from 180 videos (131 falls). The pool is model-selected and used only for training.

Second annotator. An independent second annotator labelled 100 of the test videos without seeing the first annotator's answers. On the 98 clips both labelled *fall* or *not a fall*, raw agreement was 0.796 and Cohen's $\kappa = 0.498$ [0.297, 0.678] (moderate). Disagreements were asymmetric: in 16 of 20 cases the second annotator saw a fall where the first did not (Section 5.8).

4 Experimental setup

Training data. *Staged*: UR Fall [11]; fall clips are centred on the frames labelled "falling", negatives are pre-fall clips and 2 s windows (1 s stride) over the activity videos; 316 clips, 60 falls (30 falls $\times$ 2 cameras). *Synthetic*: OmniFall OF-Syn [13]; one clip per video, centred on the fall segment for all 2,136 videos containing a fall and on the longest segment for 3,000 hash-sampled videos without one. *Real*: the 268 labelled OOPS! training clips.

External real-world test: HMDB51. HMDB51 [9] is independent of OOPS!. From its split 1 (100 clips per class) we use all 100 human-labelled `fall_floor` clips and 20 hash-sampled clips from each of the other 50 classes. A hard subset, fixed in advance, keeps falls and whole-body or ground-contact actions (cartwheel, flic-flac, somersault, handstand, dive, jump, sit, sit-up, push-up, stand, climb; 320 clips). Clips are covered by 2 s windows every 1 s and scored by their maximum window. Nothing is trained on HMDB51.

Representations and detectors. *CLIP appearance*: mean of 8 L2-normalised CLIP ViT-L/14 frame embeddings. *CLIP temporal*: [mean, last-third minus first-third, mean absolute frame-to-frame change]. *VideoMAE*: pooled VideoMAE-B (Kinetics-400) embedding of 16 frames. *CLIP + VideoMAE*: concatenation of the last two. Detectors are zero-shot CLIP (fall vs. non-fall prompt margin, prompts fixed in advance) and L2-regularised logistic regression on standardised features ($C = 0.1$, class-balanced).

Table 1: AUROC on **FallOrNot** (255 clips, 102 falls) by training data. Zero-shot CLIP: 0.717. Best per row in bold. Runs 3–4, 10.

Representation	Staged	Synthetic	Real	Syn. + real
CLIP appearance	0.723	0.728	0.732	**0.800**
CLIP temporal	0.728	0.795	**0.866**	0.823
VideoMAE	0.765	0.789	0.787	**0.842**
CLIP + VideoMAE	0.759	0.814	**0.899**	0.872

CLIP + VideoMAE 95% CIs: staged [0.698, 0.814]; synthetic [0.753, 0.869]; real [0.858, 0.935]; synthetic + real [0.825, 0.914].

Fine-tuning. VideoMAE-B fine-tuned end to end with hyper-parameters fixed before any run: 5 epochs, batch 8, AdamW, learning rate 5×10^{-5}, weight decay 0.05, one-cycle schedule, class-balanced sampling, horizontal flips, last checkpoint; trained on synthetic (1,000 + 1,000 clips), real (268 clips) or both; three seeds.

Metrics and statistics. AUROC, average precision and recall at 5% false-positive rate, with 95% confidence intervals from a bootstrap (1,000 resamples). OOPS! videos are cut from YouTube compilations, and clips from one compilation are correlated, so on FallOrNot we resample whole compilations (see the leakage check in Section 5.2); on HMDB51 we resample clips. Two detectors are compared on the same clips with a paired bootstrap of the same kind (2,000 resamples), which gives the difference, its 95% CI and a two-sided p-value. p-values are not adjusted for multiple comparisons; apart from the two tests specified in advance (fine-tuning and source-balanced mixing), all comparisons are exploratory.

Test specified in advance. Before running any fine-tuning we committed one success rule to the project repository together with the fixed hyper-parameters: fine-tuning on synthetic + real data must beat the best frozen probe on *both* test sets, with the paired 95% CI above zero. It was not registered with an external registry.

5 Results

5.1 Staged training does not transfer

All staged-trained probes reach AUROC ≥ 0.998 in five-fold cross-validation on UR Fall (grouped by sequence), and zero-shot CLIP reaches 0.912. On FallOrNot the same probes reach 0.723–0.765 (Table 1), and the best catches 30% of falls at 5% false positives (VideoMAE, recall 0.304). In-domain staged numbers cannot rank methods: every probe is at ceiling.

5.2 On FallOrNot, a few real falls beat thousands of synthetic ones

With CLIP + VideoMAE, real training beats staged training by $+0.139$ [$+0.085$, $+0.193$] ($p < 0.001$) and synthetic training by $+0.085$ [$+0.030$, $+0.137$] ($p = 0.001$). Synthetic training scores higher than staged, the direction OmniFall reports [13], but the difference is not significant here ($+0.054$ [-0.000, $+0.110$], $p = 0.053$; CLIP temporal $+0.067$, $p = 0.030$). Adding synthetic data to real data does not help this representation (-0.027 [-0.068, $+0.012$], $p = 0.19$). Recall at 5% false positives rises from 0.167 (staged) to 0.431 (real) and 0.480 (synthetic + real).

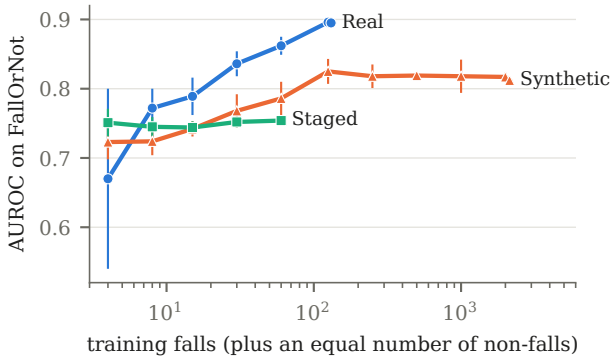


Figure 1: Equal training budget, CLIP + VideoMAE probes, FallOrNot test set (mean ± s.d. over random draws). Run 4.

Motion is what benefits from real data. Real training helps CLIP temporal ($+0.138$ [$+0.073$, $+0.198$] over staged, $p < 0.001$) but not CLIP appearance ($+0.009$, $p = 0.76$) or VideoMAE alone ($+0.022$, $p = 0.49$). Under staged training, the encoder choice makes no significant difference on FallOrNot (all $p \geq 0.30$); here the training data matters more than the encoder. This does not hold on HMDB51, where staged-trained VideoMAE is far weaker than the CLIP-based probes (Table 3).

Equal budget (Figure 1): training on k falls plus k non-falls from one source, synthetic data plateaus near 0.82 from about 125 falls up to all 2,136 (0.812–0.825), staged data stays near 0.75, and real data keeps rising: 30 real falls (0.836 $\pm$ 0.018) already exceed every synthetic budget, and about 8 real falls match the whole staged set. With only 4 falls, real data is unreliable (0.670 $\pm$ 0.130).

Falls vs. other accidents. Restricted to accident-moment clips (165 clips, 102 falls), CLIP + VideoMAE reaches 0.715 (staged), 0.772 (synthetic), 0.860 (real) and 0.848 (synthetic + real); zero-shot CLIP reaches 0.686. Telling a fall from another accident remains the hardest part of the task.

Compilation leakage check. Of the 268 real training clips, 111 come from a compilation that also contributes test clips. Retraining on the 157 clips (77 falls) from compilations absent from the test set gives 0.866, against 0.874 ± 0.014 for random subsets of the same size and class counts (CLIP temporal: 0.838 vs. 0.853 ± 0.013). We find no clear evidence that shared compilations inflate the real-data result: most of the drop from 0.899 comes from the smaller training set, and the compilation-disjoint probe still exceeds synthetic (0.814)

Table 2: AUROC on the **OF-ItW test split** (1,671 clips, 636 falls; labels by OmniFall's team). "Real (ours)": our 248 training clips; "Real (OF)": OF-ItW train + val, 734 clips. Zero-shot CLIP: 0.731. Run 11.

Representation	Staged	Syn.	Real (ours)	Real (OF)	Syn. + real (OF)
CLIP appearance	0.746	0.727	0.781	**0.829**	0.825
CLIP temporal	0.785	0.846	0.900	**0.925**	0.883
VideoMAE	0.771	0.817	0.844	**0.909**	0.899
CLIP + VideoMAE	0.815	0.872	0.927	**0.951**	0.929

and staged (0.759) training.

5.3 A third, larger test set with independent labels

To check the FallOrNot result at larger scale and with labels that are not ours, we score the same probes on OmniFall's in-the-wild test split (OF-ItW [13]): 571 OOPS! videos annotated by OmniFall's team, from which we cut 1,671 clips (636 falls; one clip per fall segment, and one per other segment of at least 1 s that is not "fallen" and does not overlap a fall). Our 20 real training clips from these videos are removed. Table 2 confirms the FallOrNot ordering, with every step significant ($p < 0.001$, paired, by compilation): synthetic over staged $+0.058$ [$+0.040$, $+0.077$], our real clips over synthetic $+0.055$ [$+0.039$, $+0.071$]. Probes trained on our 248 hand-labelled clips (0.927) come close to probes trained on OF-ItW's own 734 training clips (0.951), and on the 28 videos labelled by both teams the labels agree on 27 ($\kappa = 0.78$; a small sample skewed towards falls). Simply pooling synthetic and real clips is significantly worse than real alone here (-0.022 [-0.031, -0.013]). OF-ItW negatives are other moments of fall videos rather than other accidents, and the footage is again OOPS!, so this is a second in-domain test, not an out-of-domain one.

5.4 On an independent source, the ranking flips

On HMDB51 (Table 3, Figure 2) the real-data advantage disappears. With CLIP + VideoMAE, real training is slightly *worse* than synthetic (-0.029 [-0.057, -0.003], $p = 0.029$) and no better than staged (-0.011, $p = 0.57$). Synthetic training rescues VideoMAE ($+0.214$ over staged, $p < 0.001$), and CLIP appearance trained on real OOPS! data falls to 0.626. HMDB51 falls look closer to staged falls than OOPS! accidents do: staged probes reach 0.899 there but 0.759 on FallOrNot, and zero-shot CLIP 0.834 vs. 0.717. Many HMDB51 falls come from films.

5.5 Fine-tuning (test specified in advance)

The test specified in advance was **not met** for any of three seeds. On FallOrNot, fine-tuning on synthetic + real differs from the best probe by -0.011 to -0.002; on HMDB51 by $+0.028$ to $+0.049$ (CI above zero for two of three seeds). Averaged over seeds (Table 4), the flip reappears: real-trained models are best on FallOrNot and synthetic-trained models are best on HMDB51. Synthetic + real is within 0.008 of the best mean on both test sets, with the smallest seed-to-seed spread on FallOrNot (on HMDB51, synthetic-only varies least). Real-only fine-tuning varies strongly out of domain (HMDB51 hard subset 0.751–0.872).

Table 3: AUROC on **HMDB51** (1,100 clips, 100 falls; nothing trained on it). Zero-shot CLIP: 0.834. Run 5.

Representation	Staged	Synthetic	Real	Syn. + real
CLIP appearance	0.783	**0.813**	0.626	0.803
CLIP temporal	**0.887**	0.866	0.825	0.842
VideoMAE	0.664	0.877	0.779	**0.901**
CLIP + VideoMAE	0.899	**0.917**	0.888	0.901

Hard subset (320 clips), CLIP + VideoMAE: staged 0.833, synthetic 0.880, real 0.777, synthetic + real 0.855.

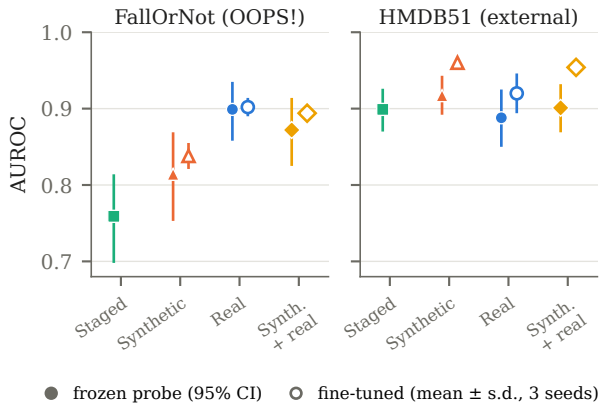


Figure 2: The ranking of training sources flips between real-world test sets. Filled: CLIP + VideoMAE probes with 95% CIs (runs 4–5). Hollow: fine-tuned VideoMAE, mean ± s.d. over 3 seeds (run 7).

5.6 Adapting with a handful of real clips: the anchored probe

The curves in Figure 1 leave a gap: with fewer than about 30 real falls a real-only probe is unstable and worse than synthetic training, and pooling the two sources lets thousands of synthetic clips outweigh the real ones (Section 5.3). We therefore test a simple alternative, a form of biased regularisation [17, 10, 12]. Let w_s, b_s be the synthetic-trained probe on features standardised with synthetic statistics. On k real falls and k real non-falls we fit $w = w_s + v$, $b = b_s + c$ by minimising

$$C \sum_i \alpha_i \, \ell\big(y_i,\ x_i^\top (w_s + v) + b_s + c\big) + \tfrac{1}{2}\|v\|^2,$$

Table 4: Fine-tuned VideoMAE, AUROC mean ± s.d. over 3 seeds. Runs 6–7.

Trained on	FallOrNot	HMDB51	HMDB51 hard
Synthetic	0.838 ± 0.017	**0.960 ± 0.006**	**0.922 ± 0.015**
Real	**0.902 ± 0.012**	0.920 ± 0.026	0.795 ± 0.067
Synthetic + real	0.894 ± 0.005	0.954 ± 0.011	0.890 ± 0.027

Table 5: Anchored probe: AUROC, mean ± s.d. over ten draws of k real falls + k real non-falls (CLIP + VideoMAE). Synthetic-only: 0.814 (FallOrNot), 0.873 (OF-ItW). Run 12.

	FallOrNot		OF-ItW test	
k	Anchored	Real-only	Anchored	Real-only
4	**.813** ± .005	.687 ± .124	**.876** ± .005	.794 ± .046
8	**.819** ± .012	.756 ± .063	**.877** ± .006	.825 ± .024
15	**.834** ± .015	.791 ± .022	**.891** ± .012	.847 ± .024
30	**.853** ± .014	.826 ± .030	**.907** ± .019	.891 ± .014
60	**.875** ± .010	.871 ± .010	**.921** ± .008	.908 ± .016
125	**.893** ± .005	.892 ± .005	.935 ± .007	**.937** ± .007
Mean	**.848**	.804	**.901**	.867

where ℓ is the logistic loss and α_i balances the classes. The penalty pulls the probe towards the synthetic solution rather than towards zero. C is chosen from $\{10^{-3}, \ldots, 10\}$ by video-grouped cross-validation on the $2k$ training clips alone.

Rule fixed in advance. Before running the method on real data we committed this rule to the repository: on both FallOrNot and the OF-ItW test split, (i) at every budget $k \in \{4, 8, 15, 30, 60, 125\}$ the anchored probe's mean AUROC over ten random draws is at least the better of the real-only mean and the synthetic-only probe, minus 0.005, and (ii) averaged over budgets it is higher than both. **The rule was met on both test sets** (Table 5, Figure 3). With four real falls the anchored probe equals the synthetic probe and is far more stable than real-only training (FallOrNot: 0.813 ± 0.005 vs. 0.687 ± 0.124); cross-validation then selects larger C as k grows, and by $k = 60$ it tracks the real-only probe.

Where it does not help. The gain is confined to small budgets. With the full real pool the anchored probe equals real-only on FallOrNot (0.899 vs. 0.899, $p = 0.96$) and is slightly but significantly worse on OF-ItW (0.943 vs. 0.951; -0.008 $[-0.012, -0.003]$, $p = 0.002$). On HMDB51, which was not part of the rule, it reaches 0.902, between real-only (0.888; $+0.014$, $p = 0.055$) and synthetic-only (0.917; -0.015, $p = 0.14$). We did not run plain pooling at these small budgets.

5.7 Which training source is safest?

No single source is best everywhere, so we ask how far each one falls below the best source on each test (Table 6). By

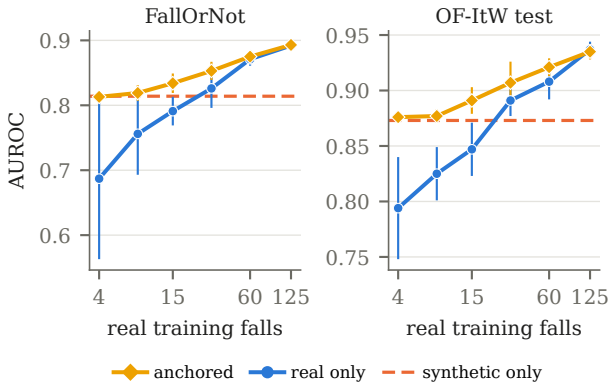


Figure 3: Anchored probe vs. real-only and synthetic-only probes as the number of real training falls grows (mean ± s.d. over ten draws). Run 12.

Table 6: Shortfall from the best training source on each test (AUROC; 0 = best). Probes: CLIP + VideoMAE. Fine-tuned: mean of 3 seeds. "hard": accident negatives only (FallOrNot) or the hard subset (HMDB51). Derived from runs 4–7.

	Trained on	FoN	FoN hard	HMDB	HMDB hard	Max
Probe	Staged	0.140	0.145	0.018	0.047	0.145
	Synthetic	0.085	0.088	0	0	0.088
	Real	0	0	0.029	0.103	0.103
	Syn. + real	0.027	0.012	0.016	0.025	**0.027**
Tuned	Synthetic	0.064	0.101	0	0	0.101
	Real	0	0	0.040	0.127	0.127
	Syn. + real	0.008	0.019	0.006	0.032	**0.032**

plain worst-case AUROC, real-only training is as good as the mixture, because HMDB51 is the easier test set. But real-only training loses up to 0.103 (probes) and 0.127 (fine-tuned) on the HMDB51 hard subset, and synthetic-only training loses up to 0.088 and 0.101 on FallOrNot's accident negatives. Synthetic + real never loses more than 0.027 (probes) or 0.032 (fine-tuned). These shortfalls are differences of point estimates and carry the uncertainty of the underlying scores.

Source-balanced mixing. In the plain mixture the 5,136 synthetic clips outnumber the 268 real clips about 19 to 1. We tested a probe in which every source-and-class cell has equal total weight, with one rule fixed in advance: it must beat the plain mixture on both test sets (paired 95% CI above zero). The rule was **not met**: balancing helps on FallOrNot ($+0.016$ $[+0.005, +0.028]$, $p = 0.004$; 0.888) but not on HMDB51 ($+0.002$ $[-0.004, +0.008]$; 0.903). The balanced probe is statistically indistinguishable from the best single source on each test set (vs. real on FallOrNot -0.011, $p = 0.52$; vs. synthetic on HMDB51 -0.014, $p = 0.09$), and its largest shortfall (0.029, HMDB51 hard subset) matches the plain mixture's.

5.8 Robustness to the annotator

Re-grading FallOrNot with the second annotator's labels (99 videos, 152 clips, 78 falls; training labels unchanged; video-level resampling) lowers absolute AUROC (best probe 0.797 instead of 0.899) but preserves the main findings: the staged-to-real gap; real over staged training ($+0.083$ [$+0.010$, $+0.157$], $p = 0.024$; CLIP temporal $+0.100$, $p = 0.010$); the learning-curve shape (about 8 real falls match staged training, 15 exceed it); and the synthetic plateau below real data at equal budget. The in-domain advantage of real over synthetic training shrinks to $+0.065$ [-0.006, $+0.137$] and is no longer significant ($p = 0.079$).

6 Discussion

Which data should a practitioner collect? In the domain where real falls are collected, a few dozen of them are worth more than thousands of synthetic falls. But that advantage does not travel: on another real-world source, synthetic data generalised better. Training on both had the smallest worst-case shortfall from the best source (at most 0.03 AUROC across all four tests, for both frozen probes and fine-tuned models), so it is the safer default when the deployment footage is not known in advance.

With few real clips, anchor rather than pool. When only a handful of real falls from the deployment domain exist, starting from the synthetic-trained probe and letting the real clips move it gave the benefit of both sources at every budget we tested, with no clear cost; once a hundred or more real falls are available, plain real-only training is as good or marginally better.

One real-world benchmark is not enough. The ranking of training sources reversed between FallOrNot and HMDB51, and the two benchmarks disagree even about how hard real falls are (staged probes: 0.759 vs. 0.899). Claims such as "synthetic beats staged" or "real beats synthetic" should be tested on more than one in-the-wild source.

False alarms on other accidents. Most fall benchmarks contain falls and everyday activities. FallOrNot's accident negatives show that separating a fall from a collision or a near-fall is the hardest sub-problem, and should be part of evaluation.

7 Limitations

- **Label uncertainty.** Agreement between annotators was moderate ($\kappa = 0.50$), and absolute scores depend on the annotator; the fall definition needs sharper written guidance.
- **Scale.** FallOrNot has 165 labelled test videos and 268 real training clips, and its confidence intervals are wide; the larger OF-ItW test (Section 5.3) confirms the ordering but shares the OOPS! source.
- **In-domain real data.** Real training and test clips both come from OOPS!; the real-data advantage on FallOrNot is partly in-domain advantage.
- **Imperfect external test.** HMDB51 falls are partly from films and trimmed to the action; they are real-world video but not necessarily real accidents.
- **Population.** Neither source focuses on older adults at home, the main use case.
- **One staged dataset** (UR Fall), which is saturated in-domain.
- **Model-selected training pool.** The real training clips were pre-selected by a CLIP zero-shot score; the test sample was not.
- **Fixed fine-tuning hyper-parameters** were not tuned; three seeds only.
- **Anchored probe.** The technique is a known form of biased regularisation; its test used frozen features, linear probes and two OOPS!-based test sets, and it gave no gain with the full real pool.
- **Many unadjusted comparisons.** Only three tests were specified in advance (one was met); the remaining p-values are exploratory.

8 Conclusion

Staged-trained fall detectors that look perfect in the lab fail on real accidents. On the source where real falls are collected, a few dozen of them outperform thousands of synthetic falls, but on an independent real-world source the ranking flips. Training on synthetic and real data together had the smallest worst-case shortfall, and with only a handful of real falls a probe anchored to the synthetic solution was never meaningfully worse than either single source and better than both on average, under a rule fixed in advance. Fall detection should be evaluated on more than one real-world source. We release FallOrNot's labels, timestamps and code.

Data and ethics. OOPS! videos are public web videos; we release only labels, timestamps and code. UR Fall is licensed CC BY-NC-SA 4.0 for non-commercial research; HMDB51 is CC BY 4.0; OmniFall annotations are released for non-commercial use.

References

[1] E. Auvinet, C. Rougier, J. Meunier, A. St-Arnaud, and J. Rousseau. Multiple cameras fall data set. Technical Report 1350, DIRO, Université de Montréal, 2010.

[2] F. Bagalà et al. Evaluation of accelerometer-based fall detection algorithms on real-world falls. *PLoS ONE*, 7(5):e37062, 2012.

[3] I. Charfi, J. Miteran, J. Dubois, M. Atri, and R. Tourki. Optimized spatio-temporal descriptors for real-time fall detection: comparison of support vector machine and Adaboost-based classification. *Journal of Electronic Imaging*, 22(4):041106, 2013.

[4] D. Epstein, B. Chen, and C. Vondrick. Oops! Predicting unintentional action in video. In *CVPR*, 2020.

[5] J. Fleming and C. Brayne. Inability to get up after falling, subsequent time on floor, and summoning help: prospective cohort study in people over 90. *BMJ*, 337:a2227, 2008.

[6] S. Y. Gadre et al. DataComp: In search of the next generation of multimodal datasets. In *NeurIPS Datasets and Benchmarks Track*, 2023.

[7] W. Kay et al. The Kinetics human action video dataset. arXiv:1705.06950, 2017.

[8] J. Klenk et al. The FARSEEING real-world fall repository: a large-scale collaborative database to collect and share sensor signals from real-world falls. *European Review of Aging and Physical Activity*, 13:8, 2016.

[9] H. Kuehne, H. Jhuang, E. Garrote, T. Poggio, and T. Serre. HMDB: A large video database for human motion recognition. In *ICCV*, pages 2556–2563, 2011.

[10] I. Kuzborskij and F. Orabona. Stability and hypothesis transfer learning. In *ICML*, PMLR 28(3):942–950, 2013.

[11] B. Kwolek and M. Kepski. Human fall detection on embedded platform using depth maps and wireless accelerometer. *Computer Methods and Programs in Biomedicine*, 117(3):489–501, 2014.

[12] X. Li, Y. Grandvalet, and F. Davoine. Explicit inductive bias for transfer learning with convolutional networks. In *ICML*, PMLR 80:2825–2834, 2018.

[13] D. Schneider et al. OmniFall: From staged through synthetic to wild, a unified multi-domain dataset for robust fall detection. arXiv:2505.19889, 2025 (v3, 2026).

[14] G. Ilharco, M. Wortsman, R. Wightman, C. Gordon, N. Carlini, R. Taori, A. Dave, V. Shankar, H. Namkoong, J. Miller, H. Hajishirzi, A. Farhadi, and L. Schmidt. OpenCLIP. Zenodo, 2021. doi:10.5281/zenodo.5143773.

[15] A. Radford et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.

[16] Z. Tong, Y. Song, J. Wang, and L. Wang. VideoMAE: Masked autoencoders are data-efficient learners for self-supervised video pre-training. In *NeurIPS*, 2022.

[17] J. Yang, R. Yan, and A. G. Hauptmann. Cross-domain video concept detection using adaptive SVMs. In *ACM Multimedia*, 2007. doi:10.1145/1291233.1291276.

[18] World Health Organization. Falls (fact sheet), 21 October 2026. `https://www.who.int/news-room/fact-sheets/detail/falls`, accessed November 2026.